

[BUG] With many observables `generate_shifted_tapes()` is called "unreas

<https://github.com/PennyLaneAI/pennylane/issues/2430>

ROW 25

[BUG] With many observables `generate_shifted_tapes()` is called "unreasonably often" resulting in massive performance loss #2430

New issue

Closed



cvijm opened on Apr 8, 2022 · edited by cvijm

Edits ...

Expected behavior

When taking the parameter shift hessian of a `QNode` returning the expectation value of several compatible observables, I expect the construction of shifted tapes to only take place once per gate, then the derivatives of all observables can be computed from this single set of shifted gates.

Actual behavior

In reality `parameter_shift.py:152:expval_param_shift` is called for each observable separately and `generate_shifted_tapes()` is therefore called number of parametrized gates times number of observables many times.

This results in a massive performance hit. In the 10 qubit example below, PennyLane spends over 60% of the time copying tapes, compared to just ~4% with the simulation of circuits, even on the slow `default.qubit` simulator.

Additional information

To get a nice graphical overview of what is happening run the example below with `cProfile`, e.g.. like this: `python -m cProfile -o timing test.py && gprof2dot -f pstats timing -o graph.dot && dot -Tsvg graph.dot -o graph.svg`

Source code

```
# test.py
# run with:
# python -m cProfile -o timing test.py && gprof2dot -f pstats timing -o graph.dot && dot -Tsvg graph.dot -o graph.svg
import pennylane as qml
from pennylane import numpy as np

num_wires = 10
wires = range(num_wires)
dev = qml.device('default.qubit', num_wires)

np.random.seed(42)
init_params = np.random.randn(num_wires//2)

@qml.qnode(dev, diff_method='parameter-shift')
def qnode(params):
    for par_idx, wire in enumerate(wires[:2]):
        qml.Hadamard(wire)
        qml.CRY(params[par_idx], wires=[wire, wire+1])
    return [qml.expval(qml.PauliZ(wire1) @ qml.PauliZ(wire2)) for wire1 in wires for wire2 in wires if wire1 != wire2]

print(len([0 for wire1 in wires for wire2 in wires if wire1 != wire2]))

for _ in range(4):
    print(qml.jacobian(qnode)(init_params))
```

Tracebacks

No response

System information

```
Python 3.8.8 | packaged by conda-forge | (default, Feb 20 2021, 16:22:27)
[GCC 9.3.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import pennylane as qml; qml.about()
WARNING: pip is being invoked by an old script wrapper. This will fail in a future version of pip.
Please see https://github.com/pypa/pip/issues/5599 for advice on fixing the underlying issue.
To avoid this problem you can invoke Python with '-m pip' instead of running pip directly.
```

Assignees

No one assigned

Labels

bug

Type

No type

Projects

No projects

Milestone

No milestone

Relationships

None yet

Development

No branches or pull requests

Notifications

Customize

Subscribe

You're not receiving notifications from this thread.

Participants



```

License: Apache License 2.0
Location: /home/cvjjm/src/covqstack/qcware/pennylane
Requires: numpy, scipy, networkx, retworkx, autograd, toml, appdirs, semantic-version, autoray, cachetools, pe
Required-by: pytket-pennylane, PennyLane-Qchem, PennyLane-Lightning, covvqetools
Platform info: Linux-5.10.102.1-microsoft-standard-WSL2-x86_64-with-glibc2.10
Python version: 3.8.8
Numpy version: 1.20.1
Scipy version: 1.6.1
Installed devices:
- default.gaussian (PennyLane-0.21.0.dev0)
- default.mixed (PennyLane-0.21.0.dev0)
- default.qubit (PennyLane-0.21.0.dev0)
- default.qubit.autograd (PennyLane-0.21.0.dev0)
- default.qubit.jax (PennyLane-0.21.0.dev0)
- default.qubit.tf (PennyLane-0.21.0.dev0)
- default.qubit.torch (PennyLane-0.21.0.dev0)
- pytket.pytketdevice (pytket-pennylane-0.1.0)
- lightning.qubit (PennyLane-Lightning-0.20.2)
- cov.qubit (covvqetools-0.1.1)

```

Existing GitHub issues

- ☒ I have searched existing GitHub issues to make sure the issue does not already exist.




cvjjm added bug on Apr 8, 2022




cvjjm on Apr 8, 2022

Contributor Author ...

Here is a screenshot of the relevant part of the profiler tree:





cvjjm on Apr 8, 2022

Contributor

Author



Also the amount of time spend on tape hashing seems excessive :-)



josh146 on Apr 8, 2022

Member



Wow, thanks for alerting us on this @cvjjm!

Also the amount of time spend on tape hashing seems excessive :-)

Yes, this one we are aware of and are thinking about how to fix! A solution probably involves smart usage of `lru_cache()`. We also noticed that calling of `np.round()` (as part of the caching) is also a significant overhead, so something we are looking into.

It is an interesting question, since the hashing results in a significant reduction in *quantum* resources (by reducing the amount of executions that take place), at a cost of additional classical resources. Something though that is very obvious in variational workflows!



cvjjm on Apr 8, 2022

Contributor

Author



Reducing quantum resources should always have priority imho. We want to use PL for real experiments! :-)



cvijm on Apr 8, 2022

Contributor Author ...

Reducing quantum resources should always have priority imho. We want to use PL for real experiments! :-)



mixd on Apr 8, 2022

Member ...

Hi @cvijm this is really helpful. Running your provided example takes around 40s locally. Attacking the offending function allows me to bring this down to ~27s. It may take a few days, but I'll aim to put together a general solution that helps here.



mixd mentioned this on Apr 11, 2022

🟢 [NVIDIA cuQuantum support for QAOA using "lightning.gpu" #2429](#)

🔗 [\[WIP\] Add support for LRU caching of param-shifted tapes #2436](#)



cvijm mentioned this on Apr 20, 2022

🔗 [Proposal for performance improvements in the Tape and Wires classes #2467](#)

cvijm on Apr 20, 2022

Contributor Author ...

Just a reminder from the discussion with @josh146 : We agree that while some smart caching can elevate the pain here and be very worthwhile overall, the core issue that should be addressed is that `generate_shifted_tapes()` should not be called `num_params x num_observables` many times but only `num_params` many times. Resolving this will dramatically improve performance for the test case above and generally help, whenever more than one observable is on a tape.



cvijm mentioned this on Apr 20, 2022

🔗 [Proposal for improvements in measurement and operation hashing #2468](#)

josh146 on Apr 20, 2022

Member ...

Thanks @cvijm; as previously discussed it is potentially non-trivial as to *why* this is happening, since there is no explicit for loop over observables in the parameter-shift logic 🤔

My gut feeling is that it might be related to Hamiltonian differentiation. Currently, when PL identifies that the observable coefficients might be differentiable, this function is applied:

```
pennylane/pennylane/gradients/hamiltonian\_grad.py  
Line 19 in c272db0
```

```
19 def hamiltonian_grad(tape, idx):
```

which *does* loop over observables.

So the solution could be identifying why this is applied when it isn't needed, and fixing it.



cvijm on May 19, 2022

Contributor Author ...

Here is some more statistics to highlight the severity of this:

- On PennyLane v0.19.0 executing the example above takes (most of this is startup time):

```
real    0m17.312s  
user    0m8.488s  
sys     0m4.016s
```



Microsoft Po



cvijm on May 19, 2022

Contributor Author ...

Here is some more statistics to highlight the severity of this:

- On PennyLane v0.19.0 executing the example above takes (most of this is startup time):

```
real    0m17.312s
user    0m8.488s
sys     0m4.016s
```



- On PennyLane v0.23.0 it takes:

```
real    2m31.671s
user    1m37.670s
sys     0m51.943s
```



The other open question was, if I understand [@josh146](#) correctly, why `generate_shifted_tapes()` is called so often because there should be no loop over observables.

Putting into `generate_shifted_tapes()` a `traceback.print_stack()` gives some insight into what is happening:

```
File "pl_test10.py", line 24, in <module>
    print(qml.jacobian(qnode)(init_params))
File ".../pennylane/pennylane/_grad.py", line 324, in _jacobian_function
    jac = tuple(_jacobian(func, arg)(*args, **kwargs) for arg in _argnum)
File ".../pennylane/pennylane/_grad.py", line 324, in <genexpr>
    jac = tuple(_jacobian(func, arg)(*args, **kwargs) for arg in _argnum)
File ".../autograd/autograd/wrap_util.py", line 20, in nary_f
    return unary_operator(unary_f, x, *nary_op_args, **nary_op_kwargs)
File ".../autograd/autograd/differential_operators.py", line 61, in jacobian
    return np.reshape(np.stack(grads), jacobian_shape)
File ".../autograd/autograd/numpy/numpy_wrapper.py", line 88, in stack
    arrays = [array(arr) for arr in arrays]
File ".../autograd/autograd/numpy/numpy_wrapper.py", line 88, in <listcomp>
    arrays = [array(arr) for arr in arrays]
File ".../autograd/autograd/core.py", line 14, in vjp
    def vjp(g): return backward_pass(g, end_node)
File ".../autograd/autograd/core.py", line 21, in backward_pass
    ingrads = node.vjp(outgrad[0])
File ".../autograd/autograd/core.py", line 67, in <lambda>
    return lambda g: (vjp(g),)
File ".../pennylane/pennylane/interfaces/autograd.py", line 188, in grad_fn
    vjp_tapes, processing_fn = qml.gradients.batch_vjp(
File ".../pennylane/pennylane/gradients/vjp.py", line 313, in batch_vjp
    g_tapes, fn = vjp(tape, dy, gradient_fn, gradient_kwargs)
File ".../pennylane/pennylane/gradients/vjp.py", line 182, in vjp
    gradient_tapes, fn = gradient_fn(tape, **gradient_kwargs)
File ".../pennylane/pennylane/transforms/batch_transform.py", line 330, in __call__
    return self._tape_wrapper(*targs, **kwargs)(qnode)
File ".../pennylane/pennylane/transforms/batch_transform.py", line 418, in <lambda>
    return lambda tape: self.construct(tape, *targs, **kwargs)
File ".../pennylane/pennylane/transforms/batch_transform.py", line 402, in construct
    tapes, processing_fn = self.transform_fn(tape, *args, **kwargs)
File ".../pennylane/pennylane/gradients/parameter_shift.py", line 667, in param_shift
    g_tapes, fn = expval_param_shift(tape, argnum, shifts, gradient_recipes, f0)
File ".../pennylane/pennylane/gradients/parameter_shift.py", line 240, in expval_param_shift
    g_tapes = generate_shifted_tapes(tape, idx, op_shifts, multipliers)
File ".../pennylane/pennylane/gradients/general_shift_rules.py", line 416, in generate_shifted_tapes
    traceback.print_stack(f=None, limit=None, file=None)
```



Can someone here who knows the inner guts of PL better than me make sense of this?



mxld on May 19, 2022 · edited by mxld

Edits Member ...

Hi [@cvijm](#) thanks for the above. I am currently going through this process, and there are a few things going on under the hood here with the 0.23.0 release. I plan to have a solution for some of these shortly.



josh146 on May 19, 2022 - edited by josh146

Edits Member ...

@cvjjm I'm digging into this in more detail, and I think this might be due to this issue in Autograd: <https://stackoverflow.com/questions/54488875/improve-performance-of-autograd-jacobian>

In particular, Autograd *itself* has a for-loop over the output shape of the cost function it is computing the Jacobian, which is quite inefficient. As far as I can tell, this doesn't happen in Torch, TensorFlow, or JAX.

For example, the following program is very fast, and doesn't have this for loop over observables:

```
import pennylane as qml
from pennylane import numpy as np
import jax
from jax import numpy as jnp

num_wires = 10
wires = range(num_wires)
dev = qml.device('default.qubit', num_wires)

np.random.seed(42)
init_params = np.random.randn(num_wires//2)

@qml.qnode(dev, diff_method='parameter-shift', interface="jax")
def qnode(params):
    for par_idx, wire in enumerate(wires[::2]):
        qml.Hadamard(wire)
        qml.CRY(params[par_idx], wires=[wire, wire+1])
    return [qml.expval(qml.PauliZ(wire1) @ qml.PauliZ(wire2)) for wire1 in wires for wire2 in wires if wire1 != wire2]

print(len([0 for wire1 in wires for wire2 in wires if wire1 != wire2]))

for _ in range(4):
    print(jax.jacobian(qnode)(init_params))
```

Then the question becomes: how come PennyLane v0.19 wasn't affected by this poor implementation of `autograd.jacobian`?

I think I can answer this: it is because, in PennyLane v0.19, we were attempting to hotfix this by caching the Jacobian *between* autograd calls! However, when we re-wrote PennyLane for arbitrary orders of differentiation, this hotfix became a memory leak and was removed.



josh146 on May 19, 2022

Member ...

It looks like attempts were made to improve the performance of `autograd.jacobian`, but as it is now no longer maintained, this will likely never happen: [HPS/autograd#524](https://hps.autograd#524)



josh146 on May 19, 2022 - edited by josh146

Edits Member ...

@mlxd @cvjjm here is the former hotfix, from PR #1131:

[pennylane/pennylane/interfaces/autograd.py](#)
Lines 198 to 204 in 8745c6f

```
198 # The following dictionary caches the Jacobian and Hessian matrices,
199 # so that they can be re-used for different vjp/vhp computations
200 # within the same backpropagation call.
201 # This dictionary will exist in memory when autograd.grad is called,
202 # via closure. Once autograd.grad has returned, this dictionary
203 # will no longer be in scope and the memory will be freed.
204 saved_grad_matrices = {}
```

It took advantage of closure in a really scary way.

It took advantage of closure in a really scary way.

In this particular case, my main suggestion would likely be to switch to JAX, or see if we could somehow modify `qml.jacobian` to compute the Jacobian via low-level Autograd functions in a way that is more efficient (see the code by [@twhughes](#) in the link above)



cvijm on May 20, 2022 · edited by cvijm

Edits Contributor Author ...

Looks like a perfectly valid use of closure to me :-)

A similar "caching" effect can also be used to speed up the above example (and still get the correct result). See here for a demo: [408ca1d](#)

```
pennylane/gradients/parameter_shift.py
144 144     return qml.math.stack([coeffs, mults, shifts]).T
145 145
146 146
147 + result_cache = dict()
147 148 def expval_param_shift(tape, argnum=None, shifts=None, gradient_recipes=None, f0=None):
148 149     """Generate the parameter-shift tapes and postprocessing methods required
149 150     to compute the gradient of a gate parameter with respect to an
150 151
151 152     @ -170,6 +171,13 @@ def expval_param_shift(tape, argnum=None, shifts=None, gradient_recipes=None, f0
152 153
153 154     list of generated tapes, in addition to a post-processing
154 155     function to be applied to the results of the evaluated tapes.
155 156
156 157     """
157 158
158 159     + global result_cache
159 160     + key = (tape, tuple(argnum), shifts, tuple(gradient_recipes), f0)
160 161     + print("aeaedaedaedx", key)
161 162     + if key in result_cache:
162 163     +     print("Have seen this tape before!")
163 164     +     return result_cache[key]
164 165
165 166
166 167     argnum = argnum or tape.trainable_params
167 168
168 169     gradient_tapes = []
169 170
170 171
171 172     @ -283,6 +291,8 @@ def processing_fn(results):
172 173
173 174     283 291
174 175     284 292     return qml.math.T(qml.math.stack(grads))
175 176     285 293
176 177     294 + result_cache[key] = (gradient_tapes, processing_fn)
177 178     295 +
178 179     286 296     return gradient_tapes, processing_fn
179 180
180 181
```

This is obviously not a finished solution, one certainly needs to free the `result_dict` somehow at the right time, but it works and gives the correct results in the example above without the insane number of tape shift/copy/hash operations.

Maybe something along those lines can be added to PL?

And btw: Good job in reducing the hashing time. This already runs much faster with master than it did with 0.22.0, but it is still much slower than it should be...



josh146 on May 26, 2022

Member ...

Thanks [@cvijm](#)! I'll have a look into it. The subtlety of the previous hotfix (which I don't think is possible with the higher order diff support) is that the cache became automatically out of scope after Autograd backpropagation, which meant that Python was clearing the memory.

To replicate that, we might have to dig in and work out how to add the caching to <https://github.com/PennyLaneAI/pennylane/blob/master/pennylane/interfaces/autograd.py>.

Having said that, I think rewriting `qml.jacobian` to use a mix of jvp's and vjp's -- that is more efficient than simply looping over vjp's --- is the right way to go here, and should negate the need to cache altogether





josh146 on May 26, 2022

Member ...

Thanks @cvijm! I'll have a look into it. The subtlety of the previous hotfix (which I don't think is possible with the higher order diff support) is that the cache became automatically out of scope after Autograd backpropagation, which meant that Python was clearing the memory.

To replicate that, we might have to dig in and work out how to add the caching to <https://github.com/PennyLaneAI/pennylane/blob/master/pennylane/interfaces/autograd.py>.

Having said that, I think rewriting `qml.jacobian` to use a mix of jvp's and vjp's -- that is more efficient than simply looping over vjp's --- is the right way to go here, and should negate the need to cache altogether



cvijm on Jul 4, 2022

Contributor Author ...

The solution you outlined with mixing of jvp's and vjp's sounds good! Is there any chance this could make it into one of the next releases?



antalszava on Jul 5, 2022

Contributor ...

Hi @cvijm, we have been looking into a fix to this issue in #2645. It would be using the caching approach. The PR needs a couple of final touches, but our plan is to have it in with the v0.25.0, in the next release. Would that work well?



cvijm on Jul 5, 2022

Contributor Author ...

Great! The implementation in #2645 looks very nice.



2



mxld mentioned this on Jul 12, 2022

🔗 [Cache jacobian for reuse with param-shift #2645](#)



antalszava on Jul 13, 2022

Contributor ...

Hi @cvijm, we've just merged in the fix. Let us know how it is in your use case. :)



cvijm on Jul 13, 2022 · edited by cvijm

Edits ... Contributor Author ...

Awesome! Just tested and current master now even beats the performance of v0.18.0 175.19s to 250.03s in a test that took several hours to run with v0.22 - v0.24 :-)



1



antalszava on Jul 13, 2022

Contributor ...

That's awesome to hear! 🎉 :) Our plan now is to ship v0.25.0 around mid-August, this improvement would be attached to that release.

I'm closing this issue, but it can be re-opened if there is/will be relevance still.

